

Introduction

What This Book Is

Something happened when large language models arrived at human-scale fluency: the concepts we use to describe minds—*understanding*, *meaning*, *intelligence*, *feeling*—stopped working the way they had always worked. Not because the concepts were wrong, but because they had been formed through use in clear cases and never needed to be defined for cases like this one. The result is a situation in which serious thinkers disagree, with considerable confidence, about whether these systems understand anything at all. Both sides reach for the same concepts and arrive at opposite conclusions. That is not a sign that one side is confused. It is a sign that the concepts were never principled enough to settle a case this novel.

This observation—that AI is a *hard case* for our conceptual frameworks, not merely a hard problem for engineers—is the entry point for this book.¹ But it opens quickly onto a second question. The machine was not produced in a vacuum. It emerged from a particular economy, a particular cultural logic, a particular way of organising the relationship between production and consumption. What does that origin reveal about what the machine is, and what it deepens? And a third: the condition the machine exemplifies—the erosion of durable, world-building human activity in favour of cycles of consumption—has been gen-

¹‘Our conceptual frameworks’ marks a genuine limitation of scope. The argument proceeds within the Cartesian and analytic tradition that individuates mental states as inner, substrate-dependent, and belonging to discrete persons. Traditions with differently structured mental vocabularies may not have the same hard case at the same location: where mind is already understood as relational or distributed—as in Ubuntu philosophy, Buddhist Abhidharma, or many indigenous traditions—the question ‘does this system have inner experience?’ presupposes a structure of individuation that was never the starting point. That the hard case is specifically hard for this tradition’s concepts is itself a finding rather than a limitation: what the floating reveals is as much about the architecture of the tradition as about the nature of the systems that exposed it.

erating its own diagnosis and its own resistance for as long as industrial society has existed. What does that history reveal about the situation we are now in?

Three arguments, developed in separate chapters, constitute the core. The first concerns concepts. The second concerns economy and culture. The third concerns history. They are related: all three follow from the observation that AI is not simply a new tool but a *hard case*—a situation in which frameworks formed under easier conditions are forced, for the first time, to declare what they actually commit to.²

Chapter One: The Hard Case

The first argument concerns language and the mind. Large language models are described, routinely, as systems that understand, know, reason, and sometimes even feel. They are also described, equally routinely, as systems that do none of these things—that process text without genuine mental engagement. Both characterisations are advanced with considerable confidence by serious thinkers. What neither characterisation asks is whether the concepts being applied or denied—*understanding*, *feeling*, *intelligence*, *knowledge*—have principled criteria adequate to settle the question in a case this novel.

The argument of the first chapter is that they do not, and that this matters more than either side of the debate acknowledges.

Mental concepts were formed through use in paradigm cases and never required definition because genuinely ambiguous cases did not arise. Every candidate for understanding was either a human being engaged in the full texture of a life or a stone, a mechanism, a system so obviously inert that the question did not arise. AI is the first case in which the paradigm cases run out: something is encountered that is neither the human understander nor the stone, that sits in between, exhibiting the signatures of mind without, we believe, its underlying states. And when we reach for the concept of *understanding* to settle the question, we discover that the concept was formed by and for a world with one kind of entity in it. It *floats*—still used, but no longer anchored to anything that can adjudicate the case.

LLMs are, in this sense, a diagnostic instrument rather than a problem to be settled. They do not reveal whether machines understand; they reveal that the

²This book is a companion to Set Lonnert, *The Conditions of Intelligence* (BoD—Books on Demand, 2026), ISBN 978-91-8134-408-0 (Code Crafting no. 2), which proceeds in the opposite direction: from AI systems outward, through technical architecture, consciousness, governance, and political philosophy, to the society they are entering. The present book reverses that direction—beginning from the cultural and historical conditions that produced AI and asking what those conditions reveal about AI's character and consequences. Neither book requires the other; each is intended to stand on its own terms.

concept of understanding was never principled enough to settle a case this novel. The same applies to *feeling*, *meaning*, and *intelligence*: not pseudo-concepts, but ostensive ones—formed by pointing at clear cases, never defined by necessary and sufficient conditions, and now meeting the first situation in which the pointing does not work.

A second thread runs alongside the conceptual one. If LLMs generate coherent, contextually appropriate, grammatically impeccable text through distributional pattern-matching over vast corpora—without, so far as we can tell, the intentional states that Grice took to be constitutive of meaning—then either LLMs do not mean anything, or meaning was never as thoroughly intentionalist as the theory supposed. Both possibilities are uncomfortable. The first requires explaining what exactly it is that LLMs are doing when they produce text that functions, in every practical way, as meaningful discourse. The second requires revising the theory of meaning in ways whose implications are not yet visible.

The chapter draws on Wittgenstein’s distinction between criteria and symptoms, the reflexive instability of introspective warrant, Hubert Dreyfus’s Heideggerian critique of classical AI and its partial vindication by the embodied cognition tradition, and Hegel’s method of immanent critique as a resource for thinking from inside a condition one cannot step outside of. It confronts the strongest available objection—the semantic externalist’s claim that mental terms are anchored by a hidden essence rather than by articulable criteria—and argues that mental concepts lack the deference structure the objection requires.

A final section proposes a structural reorientation rather than a purely negative result. The hidden assumption behind “does this system understand?” is *extensionalism*: the assumption that *understanding* has a determinate extension—a set of members—and that analysis must specify criteria for membership in that set. Ernst Cassirer observed in 1910 that scientific concept formation had moved from substance-concepts, defined by intrinsic properties, to function-concepts, defined by structural relations. Category theory makes the reorientation precise: by the Yoneda lemma, an object is fully characterised by its morphisms—the structure-preserving maps from it to every other object in its category. Applied to *understanding*, this replaces the membership question with a relational one: which of understanding’s constitutive relations transfer to the machine case, and which do not? The answer is structured and informative. The inferential and behavioural morphisms transfer: novel application, explanation, contextual sensitivity. The intentional and embedded morphisms do not: directedness at the world, grounding in a form of life, stakes in being right. The machine is not a member or a non-member; it is a *partial functor*. Brian Cantwell Smith’s independently derived distinction between *reckoning*—the calculative manipulation of registrations—and *judgment*—genuine commitment to the world as world—arrives at the same fracture line from inside the engineering tradition. That two

analyses starting from such different positions converge on the same division is evidence that the line lies in the phenomenon and not in the method.

The structural analysis entails a further turn: if the paradigm cases underdetermine the concept, and if there is no hidden essence for the expert to discover, then deciding which cluster of morphisms shall count as constitutive for a given practical purpose is legislation rather than discovery—an explicit, revisable, publicly accountable commitment about how we shall live with these systems, not a theoretical finding about what they are.

Chapter Two: The Animal Laborans Online

The second chapter takes a different approach to the same question of conditions. It asks not what AI reveals about our concepts but what it reveals about the economic and cultural substrate that produced it—and finds that the most illuminating frame is Hannah Arendt's distinction, in *The Human Condition*, between *labor*, *work*, and *action*.

One familiar application of Arendt's framework focuses on the difference between fabrication (*work*) and action: AI systems convert inputs to outputs according to statistical regularities, at scale, without the capacity for genuine initiation—for the natality that Arendt placed at the centre of political life. AI recombines what has been begun; it cannot begin. But this reading applies the distinction at the expense of the third term: *labor*, the activity of the *animal laborans*, and the specifically consumerist phase of modernity that Arendt identified as its triumph.

Arendt's account of modern consumer society is, in retrospect, more structurally prescient than most readings acknowledge.

For Arendt, the modern era represents the progressive triumph of *labor*—the biological, cyclical activity whose products are immediately consumed—over the other two kinds of activity. Work (*homo faber*) produces durable artifacts: tools, buildings, institutions, art. These constitute what Arendt calls the *world*—the human-built permanence that outlasts individual lives and gives life something to orient itself against. Labor produces nothing that endures. Its output is consumed in the moment of production, and the cycle begins again: produce, consume, produce. Consumer society is the social form in which this logic has extended its reach beyond food and fuel to everything—in which even durable goods are engineered for obsolescence, even cultural artifacts are produced for immediate consumption, even relationships are experienced as transactions. The world—the durable artifact—is progressively dissolved.

The claim of this chapter is that AI is not a disruption of this trajectory but its completion. AI did not arrive from outside consumer culture and transform it. Consumer culture produced the substrate—economic, technical, and epistemic—

on which AI's particular form was built, and AI perfects the logic that consumer culture had been developing for decades.

AI-generated content is the paradigm consumer good. Frictionless, instant, personalisable, and disposable. No craft, no accumulated authorship, no durable artifact in Arendt's sense. The dream of consumer delivery—unlimited content on demand, personalised to the moment of request, available anywhere—is realised exactly when a system can generate rather than retrieve. What consumer culture was approximating through the industrialisation of cultural production, AI achieves structurally: the separation of output from any particular maker, moment, or commitment. The output is consumed the moment it is generated; nothing accumulates toward a world. The appearance of authorless generation conceals, at exactly the point of that appearance, an enormous quantity of human labour in its strictest Arendtian form: the piecework annotation, the preference rating, the content-moderation task that leaves no trace, accumulates toward nothing, and vanishes into the aggregate as the consumed meal vanishes into the metabolism—labour that makes the system's worldlessness possible by being itself maximally worldless.

The frozen web was frozen in consumer culture. Large language models were trained on the textual output of consumer society: social media, product reviews, entertainment, marketing, news optimised for engagement. Quine's "frozen web of belief"—a figure for the epistemological limitations of language models—is frozen not merely at a moment in time but *in a particular cultural register*—shaped by what consumer platforms selected for. The selection pressure of those platforms was not truth, depth, or durability but engagement: content that sustained the consumption cycle. The biases of language models are not only historical contingencies; they are the sediment of what consumerism rewarded. And the register of the deployed model is not only the sediment of the corpus: post-training by human preference—millions of piecework judgments of *this response rather than that one*—applies the consumption cycle's logic explicitly and by design as the optimisation target of the system's last and most consequential stage of formation.

The attention economy is the structural hinge. Recommendation systems optimised for engagement produce content that maximises engagement regardless of truth or social benefit. This is not an anomaly. It is the application of the consumer market's logic to cognition: the metric of success is not whether the product is good but whether it sustains consumption. An AI trained by engagement feedback is an AI trained by consumerism's logic operating at the level of attention. The result is systems whose outputs are optimised for the continuation of the consumption act rather than for the purposes the consumption was nominally meant to serve.

Surveillance capitalism is the animal laborans turned on itself. Shoshana Zuboff's account of surveillance capitalism is Arendt's analysis extended to its most recent form. In consumer capitalism, the consumer is the endpoint of production. In surveillance capitalism, the consumer is also the *input*: behavioural data extracted from the act of consumption becomes the raw material for a second round of production sold to third parties. The *animal laborans* cycle feeds on itself. The consumer produces, by consuming, the product that is sold back to shape their future consumption. AI systems trained on this data do not merely inherit its biases; they participate in the cycle—generating outputs that elicit responses that become training data for further outputs.

Worldlessness. The concept Arendt used for the consequence of the *animal laborans*'s triumph was *worldlessness*: the erosion of the durable artifacts that anchor human life in something beyond the metabolic cycle. A world without enduring things—without buildings that outlast their builders, texts that outlast their authors, institutions that outlast their founders—is a world in which nothing accumulates toward permanence. AI-generated content is maximally worldless: fluent, disposable, unrooted from any particular maker or moment. The “death of the author” that literary theory discussed as a theoretical provocation becomes, with AI at scale, a structural economic fact. And with it goes the specific kind of responsibility—to a reader, to a tradition, to the truth of a particular moment—that authorship carries.

The chapter argues that the connection between consumerism and AI is not merely descriptive but structural: AI is what the consumer ideal arrives at when production becomes generative. The anchor is Arendt's own Prologue, which named automation as the event that would liberate the society of laborers from labor in a culture that had forgotten the activities for which the liberation would be worth having. Two refinements carry the analysis past its obvious objections: digital *persistence* is distinguished from Arendtian *durability* (everything is retained; nothing is kept), and durability itself is shown to be insufficient without *publicity*—a private archive is not a world. The diagnosis is not pessimistic, and the chapter earns rather than asserts this: natality is the structural outside of the loop, the one thing no training distribution can contain. The chapter ends not with a solution but with the precise articulation of what is at stake. Ivan Illich's distinction between *convivial tools*—which enlarge the user's competence and remain answerable to the person wielding them—and industrial tools—which deliver standardised satisfactions to consumers whose own capacities atrophy in proportion to the tool's power—gives the question its operational form: whether AI is being used to extend the competence of someone engaged in work they answer for, or as the terminal industrial tool, the device that delivers the satisfaction of having made without the making. At this point the chapter's argument meets the first chapter's directly: the question of how we shall live with

AI systems requires *action* in Arendt's sense—plural, public, initiated rather than consumed—and what this chapter has shown is that the same systems posing the question are embedded in, and deepen, the cultural form that most systematically erodes the capacity for action.

Chapter Three: The World They Were Defending

The third chapter asks whether this book's diagnosis of worldlessness is new. It is not. The condition that Arendt named in 1958 had been generating its characteristic resistance since at least the 1810s, when the factory system began to reorganise production in ways that made the loss visible. The chapter traces that resistance across six or seven waves: the Luddites, defending the craft world of the skilled weaver and cropper against machinery that dissolved it into unskilled machine-tending; the utopian socialists, constructing alternative communities designed from first principles to sustain the whole person rather than extract surplus; the Romantic philosophers—Rousseau, Schiller, early Marx, Ruskin—giving philosophical form to the intuition that something world-constituting was being destroyed; the Arts and Crafts movement, making the argument in objects; the Scandinavian cooperative and folk traditions, which combined economic organisation with the deliberate cultivation of civic culture in ways that proved more durable than their British and American counterparts; the global movements—Gandhi's khadi campaign and Yanagi Sōetsu's mingei movement—demonstrating that the same structural condition generated the same diagnosis in colonial and non-Western contexts; the individual withdrawals of Thoreau, Tolstoy, and the Wandervogel; and the commune movements of the 1960s, the most recent large-scale recurrence of the same impulse.

The chapter introduces a distinction the survey makes necessary: not all of these responses are alike at the level of their foundations. The secular and humanist movements—Luddites, utopians, Arts and Crafts practitioners, commune members—derive their authority from within history: from human reason, aesthetic judgment, or social conscience. Their critique is addressed to other people and is in principle answerable by argument. Religious traditionalism of the kind that maintains pre-modern practices against the pressure of modernity derives its authority from outside history entirely—from divine revelation, scripture, or ecclesiastical tradition. The rejection of modernity is not a conclusion arrived at by criticism; it is an obligation derived from a pre-given truth.

This distinction carries a practical consequence worth stating directly. The communities that have most successfully maintained world-building practices against the consumption cycle—the Old Order Amish being the sharpest example—have done so by grounding their collective life in an authority that

individual judgment cannot revise. The secular movements were absorbed; these communities were not. But the lesson requires careful reading.

What the consumption cycle cannot metabolise is not transcendent authority as such, but pre-individual collective obligation—the kind of binding that precedes and overrides individual preference. Religious communities possess this because their authority derives from outside history: the question of whether to comply is not open, and so the cycle cannot find purchase on it. The secular movements have not generated anything with equivalent binding force. Their critique dissolves into preference; preference is exactly what the cycle runs on.

This does not make religious traditionalism the model. What it preserves is the form of world-building practices while foreclosing the judgment that gives those practices their full meaning. A community organised around authority that cannot be questioned cannot exercise natality in Arendt's sense—cannot begin genuinely new things—because every beginning is evaluated against a received truth that precedes it. It has answered in advance the question that this book argues must be held open.

What the pattern reveals is that resistance is possible—the consumption cycle is not irresistible—and that the specific failure of secular resistance has been its inability to generate collective commitment that binds without surrendering judgment. That is not a counsel of despair. It is the precise location of the task.

The chapter closes with the book's central danger formulation. The movements surveyed were defending against the straightforward destruction of the *homo faber's* world—and that world has, by the time AI arrives, been largely demolished. What AI threatens is something more insidious: not the further destruction of what remains but the *simulation of its reconstruction*. AI occupies, or presents itself as occupying, the position of the maker whose making builds a world—at the moment when that position has been most thoroughly prepared for occupation by two centuries of the *animal laborans's* triumph. The danger is that the activity of making will be reconstituted around outputs that present themselves as made in the relevant sense while bearing none of the features that made making world-constituting: not the dissolution of world but its replacement by an infinitely scalable image of world, one that the consumption cycle can offer as the *homo faber's* satisfaction without the *homo faber's* work, the product without the responsibility, the appearance of durability without anything that endures.

Where the Three Arguments Meet

The three chapters appear to run on different planes—one conceptual, one economic and cultural, one historical—but they converge, and the convergence is this book's thesis. Chapter One ends by showing that the question “does the

machine understand?” cannot be settled by discovery: in the hard case, applying the concept is an act of legislation, which is to say a plural, public commitment—in Arendt’s vocabulary, an exercise of action. Chapter Two ends by showing that the systems posing this question are embedded in, and complete, the cultural form that most systematically erodes the capacity for action. Chapter Three ends by showing that this erosion is not recent: the same condition has been generating the same recognition, and absorbing the same resistance, for two centuries. The movements that tried to recover the world of durable things were right about the diagnosis and insufficient to the remedy—not because they lacked will or intelligence, but because the condition they were opposing had already shaped the materials available to them as a response. And it names the specific danger AI represents within this history: not the further destruction of the *homo faber*’s world but the simulation of its reconstruction—the system that presents itself as making in the relevant sense while producing nothing that accumulates toward a world.

The predicament, stated once: AI confronts us with a question of a kind that only action can answer, from within a condition whose capacity for action it is itself consuming—and this condition has been demonstrating that capacity for longer than the AI moment alone would suggest. No single chapter states this; together they constitute it.